Jiajun Fan

Portfolio: jiajunfan.com

Email: jiajunf3@illinois.edu
Google Scholar

Education

- **University of Illinois Urbana-Champaign** — Urbana, IL, USA
  Ph.D. of Computer Science; GPA 4.0/4.0 — *Aug. 2024 - Dec. 2027 (expected)*
  - **Research Field**: Stable and Autonomous RL Post-Training for Large Generative Models (Flow/Diffusion, Multi-modal Reasoning LLMs).
  - **Service**: Reviewer at ICML 2025–2026, ICLR 2025–2026, NeurIPS 2024–2026, COLM 2026, CVPR 2026, WACV 2027, AAAI 2025 & 2027, AISTATS 2025, KDD 2024.
  - **Relevant Course**: Machine Learning Algorithms for Large Language Models (LLMs).

- **Tsinghua University** — Beijing, China
  M.Eng. of Computer Technology; GPA 3.97/4.0, Top 1.3% — *Sept. 2021 - Jun. 2024*
  - **Service**: Reviewer at NeurIPS 2022–2023, ICML 2023–2024, ICLR 2024.
  - **Relevant Course**: Stochastic Processes (A). Big Data Systems (A^+). Digital Processing of Speech Signals (A). Data Visualization (A^+).

- **Nankai University** — Tianjin, China
  B.Eng. of Intelligent Science and Technology; GPA 93.28/100 (3.9/4.0), 1/83 — *Sept. 2017 - Jun. 2021*
  - **Honors**: Ranked 1st in major; National Scholarship twice (Top 1%).

Research Interest

I study how to train **multi-modal large generative/agentic models with reinforcement learning** in a way that is **stable** (e.g., lifelong continual learning)—avoiding the diversity and performance collapse that RL fine-tuning tends to cause—and **progressively autonomous** (e.g., recursive self-improvement, RSI), step by step taking the human out of the loop: first human data collection, then human labelling and hyper-parameter tuning, and finally hand-designed rewards. This work spans the following domains (**bold** = first / co-first author):

- **Data-Distribution Optimization** *(game agents)*: **GDI** [ICML'22], **LBC** [ICLR'23, **oral**, 5/4176]
- **Text-to-Image Generation** *(flow matching & diffusion)*: **ORW-CFM-W2** [ICLR'25], **ADRPO** [NeurIPS'25]
- **Audio Reasoning LLMs** *(multi-modal reasoning)*: **AudioRSI** [arXiv'26], **CESAR** [ICLR'26], **Read the Mind** [arXiv'26]
- **Tool-Use Agents** *(agentic RL)*: **Procedure-Aware RL** [COLM'26], **SpeechGym** [arXiv'26], BAR-RAG [ICML'26], PyRAG [arXiv'26]
- **AI for Science** *(protein & materials design)*: **ProteinZero** [TMLR'26], **EDiSon** [AI4Mat'24], GPFlow [NeurIPS'26]
- **Math Reasoning** *(token efficiency)*: BCR [ICML'26]
- **Embodied AI** *(VLA & efficient inference)*: ElegantVLA [NeurIPS'26], SP-VLA [ICLR'26], PRANCE [TPAMI'26]
- **Robot Design** *(morphology & control)*: **EDiSon** [AI4Mat'24]

Publications

† equal contribution. Full list with links: jiajunfan.com/publications.

First / co-first author

1. Zheng, Q.†, **Fan, J.**†, Cheng, C., Liu, G. Procedure-Aware Reinforcement Learning for Tool-Augmented Large Language Models. Conference on Language Modeling (**COLM 2026**).
2. Wang, Z.†, **Fan, J.**†, Guo, R., Nguyen, T., Ji, H., Liu, G. ProteinZero: Self-Improving Protein Generation via Online Reinforcement Learning. Transactions on Machine Learning Research (**TMLR 2026**).
3. **Fan, J.**, Ren, R., Li, J., Pandey, R., Shivakumar, P.G., Bulyko, I., Gandhe, A., Liu, G., Gu, Y. Incentivizing Consistent, Effective and Scalable Reasoning Capability in Audio LLMs via Reasoning Process Rewards. International Conference on Learning Representations 2026 (**ICLR 2026**).

4. **Fan, J.**, Wei, T., Cheng, C., Chen, Y., Liu, G. Adaptive Divergence Regularized Policy Optimization for Fine-tuning Generative Models. The Thirty-Ninth Annual Conference on Neural Information Processing Systems (**NeurIPS 2025**).
5. **Fan, J.**, Shen, S., Cheng, C., Chen, Y., Liang, C., Liu, G. Online Reward-Weighted Fine-Tuning of Flow Matching with Wasserstein Regularization. International Conference on Learning Representations 2025 (**ICLR 2025**).
6. **Fan, J.**, Tang, H., Przystupa, M., Phielipp, M., Miret, S., Berseth, G. Efficient Design-and-Control Automation with Reinforcement Learning and Adaptive Exploration. **NeurIPS 2024** Workshop on AI for Accelerated Materials Design (AI4Mat).
7. **Fan, J.**, Zhuang, Y., Liu, Y., Hao, J., Wang, B., Zhu, J., Wang, H., Xia, S.-T. Learnable Behavior Control: Breaking Atari Human World Records via Sample-Efficient Behavior Selection. International Conference on Learning Representations 2023 (**ICLR 2023**), **oral presentation, ranked 5/4176**.
8. **Fan, J.**, Xiao, C. Generalized Data Distribution Iteration. International Conference on Machine Learning (**ICML 2022**).
9. **Fan, J.** A Review for Deep Reinforcement Learning in Atari: Benchmarks, Challenges, and Solutions. In the proceedings of **AAAI-22** Workshop on Reinforcement Learning in Games.
10. **Fan, J.**, Xiao, C., Huang, Y. GDI: Rethinking What Makes Reinforcement Learning Different From Supervised Learning. In the proceedings of **AAAI-22** Workshop on Reinforcement Learning in Games.
11. **Fan, J.**, Ren, R., Li, J., Pandey, R., Shivakumar, P.G., Gandhe, A., Gu, Y., Liu, G., Bulyko, I. AudioRSI: Recursively Self-Improving Audio Reasoners. arXiv preprint arXiv:2610.36578, 2026.
12. Yang, B.†, **Fan, J.**†, Zhu, Z., Ma, H., Liu, G. Training-Free On-Policy Distillation. arXiv preprint arXiv:2610.22236, 2026.
13. Yang, B.†, **Fan, J.**†, Ma, H.†, Guo, R.†, Liu, G. Does Scaling Reinforcement Learning Really Require More Training? arXiv preprint arXiv:2610.01133, 2026.
14. **Fan, J.**, Li, J., Shivakumar, P.G., Huang, J.-H., Luo, Q., Maruf, M., Bulyko, I., Liu, G., Ren, R. SpeechGym: An Audio-Native Gym for Training Voice Agents via Reinforcement Learning. arXiv preprint arXiv:2608.26432, 2026.
15. **Fan, J.**, Li, J., Shivakumar, P.G., Luo, Q., Huang, J.-H., Maruf, M., et al. Can We Read the Mind of an Audio LLM? A Verbalizable, Multilingual Middle-Layer Workspace. arXiv preprint arXiv:2608.24958, 2026.
16. **Fan, J.**, Cheng, C., Shen, S., Zhou, X., Liu, G. Fine-tuning Flow Matching Generative Models with Intermediate Feedback. arXiv preprint arXiv:2510.18072, under review, 2025.
17. **Fan, J.**†, Ba, H.†, Guo, X., Hao, J. Critic PI2: Master Continuous Planning via Policy Improvement with Path Integrals and Deep Actor-Critic Reinforcement Learning. arXiv preprint arXiv:2011.06752, 2020.

Co-authored

18. Cheng, C.†, Ni, Z.†, Qu, Y.†, Chen, Y., Guo, R., **Fan, J.**, Liu, G. Variable-Length Generative Protein Design via Generalized Poisson Flow. The Fortieth Annual Conference on Neural Information Processing Systems (**NeurIPS 2026**).
19. Li, Y., Liu, H., Ji, K., Meng, Y., **Fan, J.**, Wang, Y., Qin, S., Wu, C., Xia, S.-T., Wang, Z. ElegantVLA: Learning When to Think for Efficient Vision-Language-Action Models. The Fortieth Annual Conference on Neural Information Processing Systems (**NeurIPS 2026**).
20. Sun, J., Jiang, P., Wang, S., **Fan, J.**, Wang, H., Ouyang, S., et al. Rethinking the Reranker: Boundary-Aware Evidence Selection for Robust Retrieval-Augmented Generation. International Conference on Machine Learning (**ICML 2026**).
21. Yang, B.†, Ma, H.†, **Fan, J.**, Liu, G. Batched Contextual Reinforcement. International Conference on Machine Learning (**ICML 2026**).
22. Li, Y.†, Tang, C.†, Meng, Y., **Fan, J.**, Chai, Z., Ma, X., Wang, Z., Zhu, W. PRANCE: Joint Token-Optimization and Structural Channel-Pruning for Adaptive ViT Inference. IEEE Transactions on Pattern Analysis and Machine Intelligence (**TPAMI 2026**).
23. Li, Y., Meng, Y., Sun, Z., Ji, K., Tang, C., **Fan, J.**, Ma, X., Xia, S.-T., Wang, Z., Zhu, W. SP-VLA: A Joint Model Scheduling and Token Pruning Approach for VLA Model Acceleration. International Conference on Learning Representations 2026 (**ICLR 2026**).
24. Wang, Z., **Fan, J.**, Nguyen, T., Ji, H., Liu, G. Variational Supervised Contrastive Learning. The Thirty-Ninth Annual Conference on Neural Information Processing Systems (**NeurIPS 2025**).
25. Wang, H., **Fan, J.**, Chen, Z., Li, H., Liu, W., Liu, T., Dai, Q., Wang, Y., Dong, Z., Tang, R. Optimal Transport for Treatment Effect Estimation. Conference on Neural Information Processing Systems 2023 (**NeurIPS 2023**).
26. Xiao, C., Shi, H., **Fan, J.**, Deng, S., Yin, H. CASA: Bridging the Gap between Policy Improvement and Policy Evaluation with Conflict Averse Policy Iteration. **NeurIPS 2022** Workshop on Deep Reinforcement Learning.

27. Sun, J., Zhang, H., Xu, X., Wang, H., Xie, Y., Jiang, P., Yue, Z., **Fan, J.**, et al. Ask to Distill: Learning to Synthesize Data for On-Policy Distillation. arXiv preprint arXiv:2610.27450, 2026.
28. Yang, B.†, Ma, H.†, Guo, R., **Fan, J.**, Zhu, X., Liu, G. Batched Contextual Reinforcement II: Does RLVR Need Long Rollout Windows? arXiv preprint arXiv:2610.22887, 2026.
29. Zhong, J.†, Wu, C.†, Yu, Z.†, Jiang, X., **Fan, J.**, Li, X., et al. DiscoveryRSI: Recursively Learning to Discover Valuable Hypotheses. arXiv preprint arXiv:2610.23839, 2026.
30. Pan, L., Yang, H., Chen, Y., Sun, Y., Lu, Y., Teng, Y., **Fan, J.**, Liu, G., Chu, Z., Wang, H. CleanRM: Reward Modeling from Noisy Feedback with a Few Clean Feedback for RLHF. arXiv preprint arXiv:2610.21138, 2026.
31. Yang, H., Pan, L., Dong, Y., Li, X., Lu, Y., Chen, Z., **Fan, J.**, Shao, S., Zheng, C., Wang, H. Not All Criteria Matter: Importance-Aware Criteria Weighting for Rubric-Based LLM Evaluation. arXiv preprint arXiv:2610.19915, 2026.
32. Yang, H., Zhang, Y., Pan, L., **Fan, J.**, Wang, M., Zhang, K., et al. RubricArmor: Adversarial Evolution Improves LLM-Based Rubric Generation. arXiv preprint arXiv:2610.05308, 2026.
33. Yang, B., Li, J., **Fan, J.**, Zhang, Y.E., Guo, R., Ma, H., et al. How Much Can Language Models Gain from Test-Time Computation? arXiv preprint arXiv:2610.01110, 2026.
34. Ma, H.†, Cao, S.†, **Fan, J.**, Yang, B., Liu, G. S^3: Spectral Null-Space Swap Makes Reasoning Models Efficient. arXiv preprint arXiv:2609.37976, 2026.
35. Sun, J.†, Shi, J.†, Xie, Y., Wang, S., Parekh, J.R., Jiang, P., Shi, Z., **Fan, J.**, et al. Retrieval is Cheap, Show Me the Code: Executable Multi-Hop Reasoning for Retrieval-Augmented Generation. arXiv preprint arXiv:2605.12975, 2026.
36. Yang, B., Guo, R., **Fan, J.**, Cheng, C., Liu, G. M3: High-fidelity Text-to-Image Generation via Multi-Modal, Multi-Agent and Multi-Round Visual Reasoning. arXiv preprint arXiv:2602.06166, 2026.
37. Cheng, C., Li, J., **Fan, J.**, Liu, G. α-Flow: A Unified Framework for Continuous-State Discrete Flow Matching Models. arXiv preprint arXiv:2504.10283, 2025.
38. Wang, H., Lian, J., Wu, M., Li, H., **Fan, J.**, Xu, W., Li, C., Xie, X. ConvFormer: Revisiting Transformer for Sequential User Modeling. arXiv preprint arXiv:2308.02925, 2023.
39. Wang, H., Chen, Z., **Fan, J.**, Huang, Y., Liu, W., Liu, X. Entire Space Counterfactual Learning: Tuning, Analytical Properties and Industrial Applications. arXiv preprint arXiv:2210.11039v2, 2022.
40. Xiao, C., Shi, H., **Fan, J.**, Deng, S. An Entropy Regularization Free Mechanism for Policy-based Reinforcement Learning. arXiv preprint arXiv:2106.00707, 2021.

Industry Experience

- **Amazon — AGI Foundations (Audio / Multi-modal LLM)** Seattle, WA, USA
  Applied Scientist Intern *May 2025 - Dec. 2025*
  - **Full-Stack Distributed RL Post-Training**: Sole owner of the **end-to-end pipeline** for multi-modal reasoning LLMs—distributed training infrastructure, algorithm and reward design, evaluation, and model delivery—and sole first author of the resulting publication.
  - **Infrastructure & Algorithm**: Built and scaled a **multi-node, multi-GPU GRPO training stack** for large-scale omni-modal LLMs, and designed **CESAR**, a process-reward framework that moves RLVR beyond outcome verification and **removes any dependence on human-annotated reasoning traces**.
  - **Impact**: Diagnosed and **resolved test-time inverse scaling**—the failure mode that had made chain-of-thought counterproductive in Audio LLMs—and delivered **state-of-the-art results on MMAU Test-mini**, surpassing Gemini 2.5 Pro and GPT-4o Audio, with near-human-level reasoning on MMSU. Published at **ICLR 2026** as first author.

Research Highlights

- **Self-Evolving RLHF for Flow Matching Generative Models**
  - **Role**: **1)** Introduced a **self-evolving RLHF** framework (ORW-CFM-W2) that enables flow matching models to continuously optimize through online reward feedback **without relying on human-collected datasets** or likelihood calculations. **2)** Derived a tractable Wasserstein-2 distance bound for flow matching models, providing the **first theoretical guarantee** for collapse-free policy evolution. **3)** Established a unified perspective connecting flow matching fine-tuning with traditional KL-regularized RL, enabling controllable reward-diversity trade-offs.
  - **Achievements**: **1)** Achieved state-of-the-art alignment while maintaining generation diversity through theoretically-grounded regularization. **2)** Validated the framework's effectiveness by successfully fine-tuning large-scale models like Stable Diffusion 3 across challenging tasks including spatial understanding and compositional generation. **3)** Published a paper called "Online Reward-Weighted Fine-Tuning of Flow Matching with Wasserstein Regularization" at **ICLR 2025** as first author.

- **Behavior Control in Reinforcement Learning**
  - **Role**: **1)** Introduced a unified framework called LBC to achieve behavior control in RL. **2)** Provided a unified perspective on diverse RL methods for behavior control and potential enhancements. **3)** Validated LBC's efficacy through rigorous theoretical support and extensive empirical experiments.
  - **Achievements**: **1)** Surpassed **24 human world records** and attained the pinnacle of performance among reinforcement learning algorithms across most tasks. **2)** Published a paper titled "Learnable Behavior Control: Breaking Atari Human World Records via Sample-Efficient Behavior Selection" at **ICLR 2023** with **oral presentation** as first author.

- **Sample-Efficient Reinforcement Learning**
  - **Role**: **1)** Introduced a sample-efficient Reinforcement Learning (RL) framework known as GDI, which achieved human-level performance by optimizing the data distribution of RL agents. **2)** Supported the efficacy of GDI with a robust foundation, including both theoretical proofs and an extensive array of experiments conducted in Atari. **3)** Provided a unified perspective on various RL algorithms with GDI.
  - **Achievements**: **1)** Outperformed prior SOTA method Agent57 with **500x less data** and **twice** the average performance. **2)** Published a paper called "Generalized Data Distribution Iteration" at **ICML 2022** as first author.

Patents

- Fan, J., Xiao, C. Unified framework for model-free reinforcement learning algorithms. CN112766497A.
- Fan, J. Hyperparameter tuning algorithm based on multi-arm gambling machine optimizer. CN112926629A.
- Fan, J., Xiao, C. An unbiased estimation algorithm of behavior value function. CN112926628A.
- Fan, J., Xiao, C. Policy gradient algorithm based on double robust qualification trace. CN112926735A.
- Fan, J., Xiao, C. Asynchronous multi-arm gambling machine hyperparameter optimizer based on electoral college voting mechanism. CN112949850A.
- Fan, J. Real-time multi-hyperparameter controller. CN113052252A.
- Fan, J. Hyperspace multi-coupling parameter optimizer based on multi-arm gambling machine combined with democratic voting. CN113052253A.
- Fan, J. Fast and generalizable hyperspace coupling multi-parameter nonlinear optimizer. CN113052248A.
- Fan, J. Reinforcement learning algorithm based on generalized combination strategy space. CN113052312A.

Awards

- Outstanding Graduates (1%) — Tianjin, China, 2021
- Excellent Graduation Thesis of Nankai University — Tianjin, China, 2021
- Tang Lixin Scholarship (1%) — Tianjin, China, 2021
- **National Scholarship**, Nankai University (1/83) — Tianjin, China, 2020
- Nomination for Zhou Enlai Scholarship — Tianjin, China, 2020
- **National Scholarship**, Nankai University (1/83) — Tianjin, China, 2019
- 3rd Prize, RoboCup@Home Education World Final — Sydney, Australia, 2019
- Bronze Medal, **ACM / ICPC** Asia Regional Contest — Xuzhou, China, 2019
- National 2nd Prize, National College Students Mathematical Contest in Modeling (5%) — Tianjin, China, 2018
- The First Prize Scholarship, Nankai University (2/83) — Tianjin, China, 2018

Research Experience

- **RL Post-Training for Audio Large Language Models** — Ph.D. Student
  UIUC, Urbana, IL, USA — *May 2025 - Present*
  - **Motivation**: Reasoning in Audio LLMs remains underexplored. Chain-of-thought reasoning often degrades inference performance—a phenomenon we term *test-time inverse scaling*—due to hallucinatory and inconsistent reasoning. Existing RL methods rely on outcome-only rewards or require laborious hand-crafted process rewards, failing to cultivate genuine reasoning capability.
  - **Results**: Proposed CESAR, an online RL framework using GRPO with multi-faceted reasoning process rewards that incentivize consistency, structured reasoning, and calibrated depth. Resolved test-time inverse scaling and achieved SOTA on MMAU Test-mini, outperforming Gemini 2.5 Pro and GPT-4o Audio.
  - **Publication**: "Incentivizing Consistent, Effective and Scalable Reasoning Capability in Audio LLMs via Reasoning Process Rewards" at **ICLR 2026** as first author.

- **Procedure-Aware RL for Tool-Augmented / Agentic LLMs** — Ph.D. Student
  UIUC, Urbana, IL, USA — *Oct. 2025 - Present*
  - **Motivation**: Existing RLVR methods for tool use reward the correct *set* of tool calls while ignoring their execution order, producing agents that are fragile in state-sensitive, multi-turn environments.
  - **Results**: Introduced a procedure-aware reward decomposition (NDCG-based trace reward, suffix-matching state reward, and structural error penalties) under GRPO. On BFCL-V3 the 3B model reaches performance competitive with proprietary models, with a **+87% relative multi-turn gain** over the correctness-only RL baseline, and consistent gains on API-Bank and Bamboogle.
  - **Publication**: "Procedure-Aware Reinforcement Learning for Tool-Augmented Large Language Models" at **COLM 2026** as **co-first author**.

- **RL Post-Training for Flow/Diffusion Generative Models** — Ph.D. Student
  UIUC, Urbana, IL, USA — *Aug. 2024 - May 2025*
  - **Motivation**: Large generative models like diffusion and flow matching models lack the ability to continuously improve from reward feedback without catastrophic forgetting or mode collapse. Existing fine-tuning methods either require expensive human-collected datasets or fail to maintain generation diversity.
  - **Results**: Achieved state-of-the-art alignment performance. Successfully fine-tuned large-scale models including Stable Diffusion 3 on challenging tasks such as spatial understanding and compositional generation, while preserving generation diversity.
  - **Publication**: Published "Online Reward-Weighted Fine-Tuning of Flow Matching with Wasserstein Regularization" at **ICLR 2025** as first author. Published "Adaptive Divergence Regularized Policy Optimization for Fine-tuning Generative Models" at **NeurIPS 2025** as first author.

- **Reinforcement Learning for Robotics** — Research Assistant
  Mila, Montreal, Canada — *Oct. 2023 - May 2024*
  - **Motivation**: Robot design and control are typically treated as separate problems, leading to suboptimal solutions. Jointly optimizing both in an end-to-end manner can significantly improve overall system performance.
  - **Results**: Combined the strengths of model-free RL and design-based RL, improving multiple metrics in MuJoCo control tasks. Introduced behavior control into design-based RL, improving sample efficiency.
  - **Publication**: Published at **NeurIPS 2024** Workshop AI4Mat.

- **Behavior Control in Reinforcement Learning** — Research Assistant
  Shenzhen, China — *Jun. 2021 - Oct. 2023*
  - **Motivation**: Existing RL algorithms struggle with complex tasks requiring diverse exploration strategies. A unified framework for controlling agent behavior through adaptive strategy selection was needed.
  - **Results**: Surpassed **24 human world records** in Atari games and achieved top performance among RL algorithms across most tasks, demonstrating exceptional sample efficiency.
  - **Publication**: Published "Learnable Behavior Control: Breaking Atari Human World Records via Sample-Efficient Behavior Selection" at **ICLR 2023** with **oral presentation** (ranked **5/4176**), as first author.

- **Sample-Efficient Reinforcement Learning (GDI)** — Research Assistant
  Beijing, China — *Sept. 2020 - Jun. 2021*
  - **Motivation**: State-of-the-art RL methods like Agent57 required enormous amounts of data to achieve human-level performance in Atari games. A more data-efficient paradigm was needed to make RL practical for real-world applications.
  - **Results**: Outperformed Agent57 with **500x less data** and **twice** the average performance on the Atari-57 benchmark, surpassing **22 human world records**.

- **Publication**: Published "Generalized Data Distribution Iteration" at **ICML 2022** as first author.

- **Model-based Reinforcement Learning Algorithm** Research Assistant
  Tianjin, China *Apr. 2020 - Sept. 2020*
  - **Role**: Introduced a groundbreaking method for achieving unparalleled performance on the challenging MuJoCo environment, which represents a significant breakthrough in the field of reinforcement learning.
  - **Achievements**: Presented a cutting-edge paper titled "Critic PI2: Master Continuous Planning via Policy Improvement with Path Integrals and Deep Actor-Critic Reinforcement Learning."

- **Multifunctional Home Service Robot** Research Assistant
  Tianjin, China *Sept. 2018 - Jul. 2019*
  - **Role**: **1)** Devised a comprehensive solution for the entire project, along with a finite state machine diagram. **2)** Employed cutting-edge ROS-based algorithms for autonomous navigation and RRT path planning.
  - **Achievements**: Presented the project in the 2019 RoboCup Sydney World Finals and won **3rd Prize**.

Academic Service

- Reviewer of The 41st Annual AAAI Conference on Artificial Intelligence AAAI 2027
- Reviewer of The IEEE/CVF Winter Conference on Applications of Computer Vision WACV 2027
- Reviewer of The Fortieth Annual Conference on Neural Information Processing Systems NeurIPS 2026
- Reviewer of The Third Conference on Language Modeling COLM 2026
- Reviewer of The Forty-Third International Conference on Machine Learning ICML 2026
- Reviewer of The Fourteenth International Conference on Learning Representations ICLR 2026
- Reviewer of The IEEE/CVF Conference on Computer Vision and Pattern Recognition CVPR 2026
- Reviewer of Transactions on Machine Learning Research TMLR
- Reviewer of The 28th International Conference on Artificial Intelligence and Statistics AISTATS 2025
- Reviewer of The Forty-Second International Conference on Machine Learning ICML 2025
- Reviewer of The Thirteenth International Conference on Learning Representations ICLR 2025
- Reviewer of The 39th Annual AAAI Conference on Artificial Intelligence AAAI 2025
- Reviewer of The Thirty-Eighth Annual Conference on Neural Information Processing Systems NeurIPS 2024
- Reviewer of The 30th SIGKDD Conference on Knowledge Discovery and Data Mining KDD 2024
- Reviewer of The Forty-First International Conference on Machine Learning ICML 2024
- Reviewer of The Twelfth International Conference on Learning Representations ICLR 2024
- Reviewer of The Thirty-Seventh Annual Conference on Neural Information Processing Systems NeurIPS 2023
- Reviewer of The Fortieth International Conference on Machine Learning ICML 2023
- Reviewer of The Thirty-Sixth Annual Conference on Neural Information Processing Systems NeurIPS 2022